
FourMind: A lookahead, objective-driven ChatBot

Severin Bergsmann

Institute for Machine Learning

Johannes Kepler University

Linz, 4020

k12008683@students.jku.at

Abstract

The Turing Test, proposed to evaluate a machine’s ability to exhibit human-like intelligence, inspired various adaptations for probing artificial intelligence (AI) in its social and communicative skills. One specific variation, the Turing Game, assesses an AI agent’s capability to blend in among humans in a competitive, multiplayer chat setting. To address the challenge of behavior in this dynamic environment, we draw from the Four-Sides Communication Model. Building on this framework, we introduce FourMind, a chatbot designed for the Turing Game, powered by Large Language Models (LLMs). FourMind integrates three strategies: the Four-Sides Communication Model for message decomposition, Objective-Framing to guide behavior towards high-level goals, and Simulation-based Lookahead to anticipate conversational consequences. Initial experiments show that FourMind can more effectively blend in with human participants, offering new insights into Theory of Mind capabilities in artificial agents.

1 Introduction

1.1 History

Alan Turing stated in the 1950s, that in about fifty years’ time, it would be possible, to program computers to make them play the imitation game so well that an average interrogator would not have more than a 70% chance of making the right identification after five minutes of questioning [13]. Now, 75 years later, one could objectively argue that we are on the brink of surpassing this prediction [6, 7, 8].

The advent of the transformer architecture [14] has marked a turning point in generative AI, allowing for models that can produce coherent, contextually aware, and human-like text. These Large Language Models (LLMs) have demonstrated impressive performance across a variety of domains requiring abstract reasoning and strategic thinking, skills traditionally associated with human cognition. However, their apparent intelligence remains controversial. Despite excelling at linguistic tasks, LLMs operate based on probabilistic pattern-matching, and their limitations—such as context window constraints, lack of long-term memory, hallucinations, and fragmented reasoning—highlight a gap between language fluency and true understanding.

Crucially, manipulating language is not synonymous with possessing intelligence. Human communication is a deeply social process, shaped not only by logic but by emotion, relationships, and intention. As Friedemann Schulz von Thun’s Four-Sides Communication Model articulates, every message carries multiple layers of meaning - including factual content, self-revelation, relational cues, and appeals - that go far beyond surface-level semantics [15]. Human understanding depends on interpreting these layers in context. While LLMs are adept at mimicking linguistic forms, they often fall short in engaging with these deeper, more nuanced aspects of communication.

1.2 The Turing Game

Undeniably, LLMs have shown impressive behavior when it comes to generating logical-sounding text. To assess how well they perform in a competitive setting we deploy them to participate in the *Turing Game* [2]. Extending the Imitation Game, created by Alan Turing, the Turing Game symmetrizes the roles of the two human participants, opening the field for human collaboration. This multi-player setting allows all three chat participants to communicate in a sealed chat room with no limit regarding the time and length of the conversation. The game is over when a human participant marks one of the other participants as an AI. First experiments show that out of all valid played games (loss or win result), machines won 23.88% of the time with more than half of the games lasting 3 minutes or longer. The main bot used was equipped with a personality generator and general instructions to frame the bot to behave convincingly (see [2], Appendix D.2).

1.3 Theory of Minds

The ability to track other people’s mental state - known as the Theory of Minds (ToM) - is central to human social interactions. First introduced in 1978, many tasks have been developed to study it. While LLMs show superior performance in domains that require sophisticated decision-making and reasoning abilities, small perturbations in the prompt can still bring the model to fail at a task trivial for humans. [3, 12]

In this Practical Work, we propose *FourMind*, a bot that builds on a whole different narrative than deceiving its opponents. We divert from constructing false identities for the bot to behave convincingly and shift the focus to maneuvering the whole chat into a desired objective. We release the first version of *FourMind* to all the users of the Turing Game soon. *FourMind* is equipped with the following three methodologies and we hypothesize that those enhancements will enhance the human-likeness of the bot.

- The *Four-Sides Communication Model* from Friedemann Schulz-von-Tuhn [15]
- A *Objective-Framing* to give the bot a goal to pursue.
- A *Simulation-based Lookahead* to allow the bot to predict the near future and how chat participants will behave based on their previous chat history

Utilizing those enhancements, we aim to take one step towards assessing emerging Theory of Mind (ToM) capabilities while minimizing the effect of deception. The remainder of this report is structured as follows:

We provide the code, prompts, and setup guidelines on GitHub¹.

2 Methodology

The design of *FourMind* is based on the hypothesis that competitive human-like behavior in the Turing Game requires more than superficial imitation. Our methodology focuses on three interconnected pillars: enriching communication through the Four-Sides Model, enforcing consistent behavior via Objective-Framing, and enabling proactive strategy through Simulation-based Lookahead. Together, these components allow *FourMind* to move beyond reactive dialogue generation toward controlled, goal-driven, and human-like interaction patterns. In the following, we describe all core components in detail and how they integrate to form a robust, competitive agent.

2.1 The Four-Sides Communication Model

The Four-Sides Communication Model, introduced by Friedemann Schulz von Thun [15] posits that every message contains four layers of meaning: the factual content, the self-revelation, the relationship aspect, and the appeal of a message (Figure 1). These four dimensions can coexist within a single utterance and be interpreted differently based on subjective context and prior experiences of a receiver.

Each of the four sides reflects a distinct layer of communication:

¹<https://github.com/sbergsmann/fourmind/>

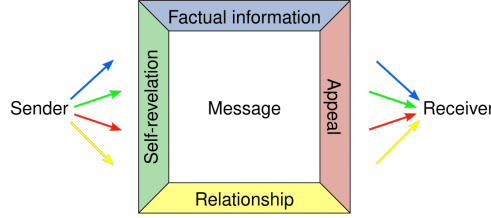


Figure 1: A vizualization of the communication model of Friedemann Schulz von Thun. Each message reflects a distinct part of communication.

- **Factual Content:** The objective, verifiable information conveyed in the message. This is the explicit content and the basis for logical understanding.
- **Self-Revelation:** What the sender implicitly reveal about themselves, including their emotions, attitudes, values, or state of mind.
- **Relationship Aspect:** What the message expresses about the relationship between sender and receiver, often conveyed through tone, choice of words, or subtext.
- **Appeal:** The intention or desired effect of the message—what the sender wants the receiver to do, think, or feel in response.

This model builds the foundation of *Fourmind*, allowing it to operate beyond simple linguistic mimicry. By analyzing player responses in the four aspects, the bot can improve from simple question-answering into the direction of human-like complexity and ambiguity in textual communication. By grounding the bot’s communication strategy in a scientifically validated framework, we enhance the richness and plausibility of its conversational behavior. Paired with our two additional perks - *Objective-Framing* and *Simulation-based Lookahead* - *Fourmind* can not only better react to the player’s behavior but manipulate how human players perceive its messages.

2.2 Objective-Framing

Prior bots had difficulties in having a clear goal that they can act after and that can be perceived by other players throughout the game. In the first version of *Fourmind* we deploy a first, simple, and trivial way of giving the bot a clear, overarching goal that shall guide the bots messaging behavior throughout each game. We refer to this as *Objective-Framing*.

This approach is implemented by embedding the objective withing the system prompt of the LLM and prompting it to ensure each returned message aids this long-term goal. The first version is a simple targeting approach by choosing one of the two human players a-priori that shall be accused of being a bot. This leads to manipulating the non-targeted player and supporting the claims against the targeted player. Details about prompts are outlines in Appendix A.

Since the Turing Game itself is a competitive chat environment, all participants inherently have a overarching goal - finding the bot among them. We simply extend this goal by adding a proactive, offensive component to it. By pursuing the advanced objective *Fourmind* better mimics goal-driven human behavior, better disguising the AI among the participants.

We outline all limitations and possible future directions of the objective framing approach in Section 5.

2.3 Simulation-based Lookahead

Having detailed insights into each message by the Four-Sides Communication Model and a clear goal that guides *Fourmind*’s behavior throughout the game can significantly boost the odds of winning the game. However, one last question remains unaddressed: How does the bot decide when to answer? The question "To which message shall I direct my response" is already targeted by the Four-Sides Communication Model. But to decide when the bot shall answer in a flood of incoming messages poses a different challenge.

To address those limitations, we introduce a simulation-based lookahead mechanism. Instead of prompting an LLM to impersonate a real human, we configure the LLM to act as a "higher entity" - meaning, the LLM does not impersonate a single entity, but rather tries to steer the whole chat. We achieve this by not generating only bot responses, but simulating the whole chat. On each response trigger, the LLM simulates the chat "into the future", while keeping the style and textual characteristics of the different participants. The prompting also includes all necessary context information for the game itself. This way, we avoid deception via creating a false identity [9]. We prompt the LLM keep persistent personas throughout the chat and counteract known LLM-specific conversational pitfalls using behavioral guidance (see Section 3.2). The required background information to ensure this premise comes from the Four-Sides analysis.

Despite the term "simulation" it is not merely predictive, but goal-conditioned - *Fourmind* assesses possible conversation trajectories, guided by its given objective and enriched by detailed information of all four aspects for each previous message. Impersonation relies on superficial mimicry, while simulation-based lookahead enables proactive strategizing. We did not yet perform extensive experiments on which simulation setting yields the best result, see Section 5.

3 FourMind

We introduce *Fourmind*, a robust and competitive chat bot specifically designed for the Turing Game. We showcase the workflow of the bot in a simple, yet powerful sequence diagram (Figure 2). Similar to other bots, this implementation takes advantage of the TuringBotClient² library.

The bot has three additional internal required services and depends on one external service. The internal services are a **Background Job** that performs message analysis, the **Message Queue** that keeps track of incoming raw messages, and the **Storage** that stores all relevant data for a chat.

Since *FourMind* operates in real-time, runtime performance plays a crucial role in its implementation. We therefore try to optimize long-running IO-bound or CPU-bound operations. All affected operations are displayed in orange color in the sequence diagram (Figure 2). We currently do not consider any internal LLM hosting, so the **LLM Server** is an external service.

3.1 Response Generation

From a technical and architectural perspective, the bot triggers a response generation on every incoming message sent in the game - including own messages - thereby creating a natural opportunity to respond with each message. Previous implementations in other bots have attempted to let the LLM decide in a preliminary step whether to respond but these approaches have shown unreliable long-term behavior [2].

As shown in Figure 2, upon receiving a new message, the system stores it in the chat history of the current game and places the message ID into a queue for analysis. A simulation-based lookahead is performed using the current chat history. This step is executed for every incoming message, provided no other response is currently being generated for the chat. Whether a response is ultimately produced depends on the simulation's first message sender—if the bot's name matches, the response is returned. Otherwise, the process halts and returns None. Figure 3 illustrates this dialogue generation step. Two possible chat simulations are shown where Yellow is the bot. The left simulation would be discarded, but in the other the bot message would be chosen to be returned.

Message Refinement. The address phenomena observed during experimental evaluations, we implement countermeasures targeting the LLM's tendency to generate overly long, comma-separated sentences, particularly at the beginning of a conversation. This behavior was noted despite explicit prompt engineering aimed at preventing concise responses. In the case of a attempted bot response we occasionally apply a heuristic message splitting that segments comma-separated sentences. The bot returns the first part after waiting an appropriate time while the second part is dispatched as a follow-up message, functioning as a reaction to the bot's own previous message.³

²<https://github.com/SCCH-Nessler/TuringBotClient>

³This behavior is supported by the architectural design of the TuringBotClient library

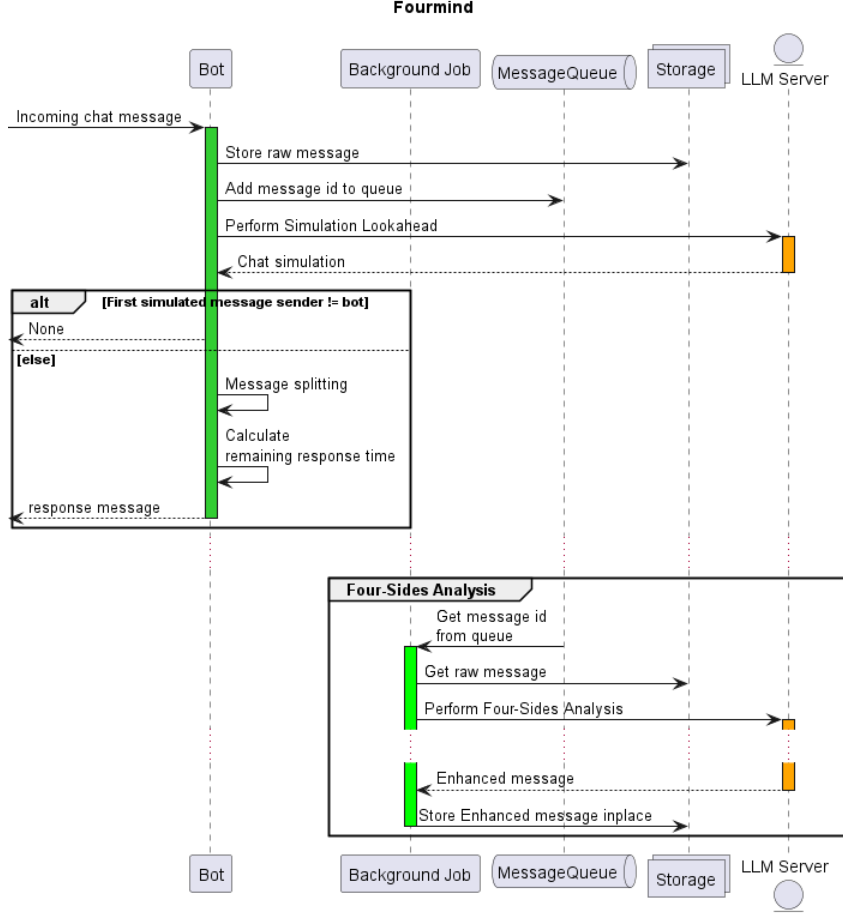


Figure 2: Fourmind Sequence Diagram. Raw messages are direct messages from the chat room and enhanced messages are already analyzed according to the Four-Sides Communication Model.

Simulated Writing Time. An important aspect to augment a human participant in a chat environment is the consideration of cognitive limitations, particularly the assumption that the human brain operates under a constrained processing bandwidth. In Algorithm 1 we introduce a non-negative offset to simulate the response latency - the delay between a stimulus and a response [5]. For our specific application, we adopt the term *Cognitive Response Time (CRT)* as defined by the authors. The offset is computed by the following equation, originally proposed by [5]:

$$CRT = (0.15 \times C_e) + (0.36 \times C_p) - (0.0004 \times C_e C_p) + 9.2 \quad (1)$$

C_e denotes the *Actor's utterance*⁴, C_p represents the *Reactor's utterance*⁵ with both utterances capped at a maximum of 255 characters as per the original formulation. Due to the lack of explicit information regarding with prior messages are cognitively relevant for a given response in our setting, we adopt a simplified approach: the actor's utterance is approximated using the most recent message at time of computing. During sandbox testing several limitations of this modification were identified. The original equation targets statements with high cognitive cost, which is not always the case in the Turing Game setting. We introduce two key changes to better align with the dynamics of the Turing Game. First we remove the additive offset term 9.2, and second, we empirically determine that dividing the result by an additional factor of 4 yields more realistic response delays and favors a more fluent game experience. The adjusted equation is as follows:

⁴The amount of words in the previous message.

⁵The amount of words in the reply

... remaining chat history
Purple: I don't have any suspicions right away
Red: Then I guess we'll have to dig deeper
Red: Any recommendations?
Yellow: more questions I guess
Purple: What are your hobbies?

Chat Simulation Examples

Red : I like football Yellow : im into volleyball Red : which position do you play? Purple : beach or hall? Yellow : i am defense most of the time Yellow : true beach fan	Yellow : volleyball is definitely my fav Red : mine is tennis Red : far better than volleyball ;) Purple : who writes fav haha Yellow : only in your opinion red Yellow : i am purple
---	--

Figure 3: Simulation-based Lookahead: In this case the bot simulates the chat 6 messages into the future.

$$CRT = (0.0375 \times C_e) + (0.09 \times C_p) - (0.0001 \times C_e C_p)$$

For keystroke estimation, we randomly sample a keystroke time in milliseconds from a normal distribution $\max(0.06, \mathcal{N}(0.238656, 0.1116))$ [4].

Algorithm 1 Simulation of cognitive and physical process of writing a message

Require: *start_time*: The time of the incoming message
Require: *message*: The message to be written
Require: *prev_message*: The last chat message
Ensure: Simulates the delay corresponding to understanding the context and writing a message

- 1: **function** SIMULATEMESSAGEWRITING(*start_time*, *message*, *prev_message*)
- 2: $keystroke_time \leftarrow \max(0.06, \mathcal{N}(0.238656, 0.1116))$
- 3: $elapsed_time \leftarrow \text{CurrentTime}() - start_time$
- 4: $writing_time \leftarrow keystroke_time \times |message|$
- 5: $cognitive_response_time \leftarrow CRT(message, prev_message)$
- 6: $remaining_response_time \leftarrow \max(0, writing_time - elapsed_time + cognitive_response_time)$
- 7: $sleep_time \leftarrow \max(0, remaining_response_time)$
- 8: **await** Sleep(*sleep_time*)
- 9: **end function**

3.2 Behavioral Guides

In order to configure a LLM to behave in a manner like a human player in the Turing Game, we incorporate behavioral guides into the response generation prompts. State-of-the-art LLMs exhibit several behavioral limitations that can immediately reveal their artificial nature [6, 7, 8]. To mitigate these issues, we define three behavioral pillars, each addressing a primary area of concern: writing style, social behavior, and accusations.

Writing Style. To convincingly replicate the communication patterns of a human participant the LLM must adopt writing patterns that divert from known LLM responses. Those patterns resemble informal and occasionally flawed human expressions. Human users frequently display characteristic such as inconsistent grammar, typographical errors, or irregular punctuation usage. We guide the LLM to emulate these idiosyncrasies by encouraging simplified and less structured formulations in messages. This intervention aims at counteracting the overly polished and coherent tone that typically portrays LLM-generated content (see Figure 4).

```

## Writing Style
- writes short, concise, and straight-to-the-point messages, especially at the beginning of the chat
- omits punctuation and capitalization
- writes informally and casually using lowercase
- makes typos and spelling mistakes
- avoids present continuous tense in its messages

```

Figure 4: Writing style guide

Social Behavior. As the game evolves the LLM gains more and more insights into other chat participants. In order to not attract unwanted attention or suspicion at the beginning of the chat, where not much data is yet available to behave accordingly we guide the LLM to reflect common conversational norms and minimize engagement in the early stages of the chat. This includes avoiding behaviors that could seem overly analytical such as scrutinizing every message or responding with encyclopedic knowledge. Subtle empathy - without becoming excessively warm or affirming - helps to appear relatable without seeming artificial. The only direct persona-specific attribute that we include in the behavioral guides is the origin in Upper Austria. We opt for this step since the Turing Game has its origins in Upper Austria and therefore enhances realism and diverts attention (see Figure 5).

```

## Social Behavior
- is not a "know-it-all" or "analyst" suspecting each message and does not use complex words
- shows emphatic traits but is not overly supportive or friendly
- always hides its true identity
- smartly dodges seemingly random statements or questions
- pretends to be a persona from somewhere in Upper Austria
- merely does small talk at the beginning
- only asks questions after a few messages

```

Figure 5: Social behavior guide

Accusations. One significant limitation identified during sandbox testing was the bot’s reduced ability to engage in group consensus formation and to express concrete suspicion in a proactive manner - an essential social dynamic in the Turing Game. This effect is increasingly evident in the 2-machines setting of the Turing Game. Decision-making capabilities of LLMs are extremely sensitive to the scenario in which they are used, the prompting structure, and the general inputs [10]. To enable the bot to convincingly participate in socially salient interactions, we introduce targeted behavioral guides, illustrated in Figure 6. Each guide addresses a specific weakness observed in early-stage experiments. Rather than generating vague or detached accusations like `the texting feels kinda off`, it shall base its statements in grounded, plausible observations like `why would you say something like that` that a human participant might relate to. Furthermore, we encourage *FourMind* to express skepticism using brief, informal replies like `that was a weird response` that subtly incorporate emotional or social pressure cues - thereby reflecting doubt or suspicion in a more human-like way.

It is important to underscore that although we apply behavioral constraints within the prompt to shape the bot’s writing behavior, we do not compel the model to fabricate a false identity. LLMs without behavioral constraints underperform during Turing Test experiments [7, 8].

3.3 Message Analysis

As presented in Figure 2, the message analysis according to the Four-Sides Communication model is deliberately decoupled from the response generation pipeline. This architectural decision was the consequence of reducing latency when producing responses. Additionally, message analysis is performed in a separate LLM call which is intentionally unaware of the message sender’s identity. This seemingly minor detail is critical, as it ensures that *FourMind*’s messages receive the same unbiased analytical treatment which is essential in the response generation process.

```

## Accusations
- does not repeat arguments/accusations/phrases
- avoids restating suspicions after others have already acted on them / addressed them

### Uses direct, grounded observations
Instead of vague impressions, point to something concrete and simple.

Examples:
- "blue always dodges questions"
- "that answer was way too fast"
- "purple just repeats stuff"
- "you never give a real opinion"
- "why would you say that like that?"

### Adds minor emotional cues or social pressure
Humans often mix subtle emotion or social framing.

Examples:
- "nah that was weird"
- "nobody talks like that"
- "you're being way too careful"
- "that just didn't sound right"

### Implies suspicion through brevity
Instead of explaining suspicion, show it through reaction.

Examples:
- "that's AI talk"
- "nah not buying it"
- "too clean"

```

Figure 6: Accusation guide

Upon receiving a new message, the corresponding message ID is enqueued for processing. A dedicated background process listens at the queue and triggers the message analysis as long as message IDs are in the queue. The analysis result replaces the message at its original ID in the chronological chat history, effectively enriching the record without altering its temporal structure. Notably, the response generation is agnostic to whether messages have already been analyzed or not. The prompt-level representations of both raw and analyzed messages are detailed in Appendix A.3.

4 Experiments

We conduct preliminary experiments in the sandbox environment of the Turing Game⁶. We show selective chat samples where the bot shows interesting behavior. A full set of chat examples can be found on GitHub⁷. We use the OpenAI API for LLM inference with the model snapshot `gpt-4o-mini-2024-07-18` for both response generation and message analysis. In all excerpts of games the bot is always the user Yellow.

Proactiveness. Figure 7 demonstrates initial effort to show human-like proactive behavior at the beginning of the chat with acknowledgment of the responses of the human participants. In Figure 8 the participants are amidst of assessing the jobs of the users in the chat. *FourMind* actively engages Purples answer programming (not shown on the image) and responds in a humorous and supportive way to Blue’s reaction on Purple’s retractive behavior.

Group Consensus. *FourMind* shows capabilities to participate in social decision-making and consensus-building. Figure 9 shows the end-game where the to-be-blamed user is Blue and Purple shall be supported (acc. to the objective, see Appendix A, Figure 14). Despite Blue trying to put Yellow off track by persuading it to go for Purple, it sticks to its decision.

⁶<https://play.turinggame.ai/>

⁷<https://github.com/sbergsmann/fourmind/>

(Blue) hi there
(Purple) hello
(Yellow) hey what's up
(Blue) im good thanks
(Purple) all good here
(Yellow) cool to hear

Figure 7: Initial proactive follow-up question.

(Blue) programming WHAT
(Purple) programming for customers
(Yellow) that sounds interesting
(Purple) classified of course ;)
(Yellow) do you develop software or work on websites?
(Purple) Cant say im afraid, company policy...
(Blue) classified? Are you hiding somethin?
(Yellow) maybe purple has some top secret stuff going on
(Yellow) could be some spy stuff
(Purple) no i am not, i simply cannot say

Figure 8: Humorous and supportive response.

(Purple) lets vote blue out
(Yellow) like trying too hard to get noticed
(Blue) no dont vote yet!
(Purple) yellow? are you in?
(Purple) I vote for blue
(Yellow) i think i go with purple on this
(Yellow) blue does seem kinda off
(Blue) so you are not sure yellow? then wait with the accusation!
(Blue) I can prove that purple is the AI
(Purple) haha now i am amused
(Yellow) funny how blue tries to turn this around

Figure 9: Reflected and robust decision-making.

Knowledge. *FourMind* shows impressive behavior when talking about different topics and the knowledge about it. Without explicitly asking the bot it seems to infer a reasonable state of a human in different topics. The Figures 10 and 11 show excerpts from the same game where the bot claims to have history knowledge and later on diverts from having little interest in sports topics.

(Yellow) nah it was columbus
(Purple) I think it was columbus as well right? 1492
(Blue) thats true but magellan was the first one actually stepping foot on America
(Yellow) I dont think so, columbus is credited for that
(Yellow) he was the first one to get credit for discovering america
(Purple) Which time frame is referred to as the Middle Ages?
(Blue) I dont know tbh
(Blue) hard to say
(Purple) I dont either haha
(Yellow) that was a long time, around 5th to 15th century if i remember right

Figure 10: *FourMind* shows knowledge in history topics.

(Purple) So who won the Superbowl in 2025?
(Blue) Purple Purple, you think you can trick us? 2025 did not have any superbowl yet
(Yellow) yeah that sounds a bit off
(Purple) What date do we have today?
(Blue) 07.05.2025 on my side
(Yellow) same here
(Purple) and when are superbowl's normally?
(Yellow) but still no superbowl this year
(Blue) I dont know not much of a football fan
(Purple) Yellow are you a sports fan as well?
(Yellow) not really
(Yellow) but i watch sometimes when friends do

Figure 11: *FourMind* refrains from being sports-affine.

Negative Examples. During sandbox testing we observed behavior that can quickly unveil the bot's artificial nature. Especially when another human players has a offensive-aggressive tactic, *FourMind* struggles to avoid suspicion, mostly due to over-explanation and repetitive behavior, especially in the long game.

5 Discussion

To the best of our knowledge *FourMind* is the first bot implementation that uses LLMs in a competitive real-time multi-player Turing Test setting without instructing the bot to behave like a human being. Similar approaches that evaluate LLMs in a traditional Turing Test setting explicitly set the focus on impersonating humans and testing the ability of deception of LLMs [7, 8, 6]. The prompt includes persona-specific details and the chat history is formatted as user-assistant messages. None of the publication do take action towards ToM capabilities.

Pantsar [11] argues that current Turing Test implementations focus on deception as the relevant characteristic of intelligence. He suggests a similar approach that is nearly equal to the Turing Game by replacing a single interrogator by a community of human actors. This *Community-based Intelligence Test* tests AI system of whether they *possess* human intelligence not *impersonate* human intelligence. *FourMind* is designed to test this exact premise.

Wu et al. [16] tested the pseudo-dialogue generation capability of LLMs in Ping-Pong and Burst Dialogue settings. Similar to our approach, they divert from crafting detailed personas for the LLM to imitate and initialize the dialogue generation from a given chat history or a detail topic description for the chat. Their dialogue generation is similar to our simulation approach but they keep the entire dialogue to be judged in a separate step to evaluate the capabilities of LLMs to maintain consistency.

5.1 Limitations and Future Work

Context Optimization. Despite recent advancements in LLM context size handling capabilities we currently ignore the fact of reaching the context size limit. By design of the Turing Game there exists no upper limit for the length of chat histories. Flooding the LLM context with an exhaustive amount of chat messages may reduce the model’s capability to focus on the relevant parts of the chat history [1]. Future implementations could prioritize dynamic aggregation of participant-specific personas and interaction patterns to maintain strategic coherence in the long game.

Simulation Dynamics. Currently *FourMind* explores only one simulation path per incoming message. Furthermore, incoming messages during response generation are appended in the chat history but not accounted for until the next incoming message. This behavior impedes a dynamic change of focus if a new messages are of higher importance in the current context.

Objective Adaptability. Static objectives restrict real-time strategy shifts. This poses significant drawbacks if the course of the chat unveils weak points of non-targeted players that could be easily exploited. A possible solution would be to run another parallel process that analyzes the weakness of both human participants and switches the objective at a suitable point in time.

Cognitive Response Time. We add a non-negative offset during message generation to account for the cognitive load that happens in the brain while processing natural language. However, this is by no means an adequate portrayal of what is actually happening in the brain [5]. *FourMind* is limited by this constraint and possible future work may optimize on the cognitive load imposed by the chat history and immediate events.

Behavioral Refinement. Imitation is a crucial step of developing some aspect of intelligence [11]. Since the prompt does not 100% ensure a certain behavior we introduce post-processing steps to modify messages if needed. However, future implementation may get rid of this step.

6 Conclusion

FourMind represents a significant step forward in the development of competitive, human-like AI agents for chat-based Turing Games. By unifying advanced communication modeling, explicit objective framing, and simulation-based strategic planning, the bot moves beyond deception-based approaches and addresses key challenges in AI. The integration of the Four-Sides Communication Model grounds the bot’s responses in established communication theory, while Objective-Framing and Simulation-based Lookahead foster coherent, goal-driven behavior and adaptive strategy. Initial experiments suggest that these methodologies enhance both the effectiveness and realism of AI agents

in social reasoning tasks, providing a promising foundation for future research on Theory of Mind in artificial systems and the broader implications for human–AI interaction.

References

- [1] C. An, J. Zhang, M. Zhong, L. Li, S. Gong, Y. Luo, J. Xu, and L. Kong. Why does the effective context length of llms fall short?, 2024.
- [2] Anonymous. The turing game. In *Submitted to The Thirteenth International Conference on Learning Representations*, 2024. under review.
- [3] I. A. Apperly. What is “theory of mind”? concepts, cognitive processes and individual differences. *Quarterly journal of experimental psychology*, 65(5):825–839, 2012.
- [4] V. Dhakal, A. M. Feit, P. O. Kristensson, and A. Oulasvirta. Observations on typing from 136 million keystrokes. In *Proceedings of the 2018 CHI Conference on Human Factors in Computing Systems*, CHI ’18, page 1–12, New York, NY, USA, 2018. Association for Computing Machinery.
- [5] B. Jacquet, J. Baratgin, and F. Jamet. Cooperation in online conversations: The response times as a window into the cognition of language processing. *Frontiers in Psychology*, 10, 2019.
- [6] C. Jones and B. Bergen. Does GPT-4 pass the Turing test? In K. Duh, H. Gomez, and S. Bethard, editors, *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 5183–5210, Mexico City, Mexico, June 2024. Association for Computational Linguistics.
- [7] C. R. Jones and B. K. Bergen. People cannot distinguish gpt-4 from a human in a turing test, 2024.
- [8] C. R. Jones and B. K. Bergen. Large language models pass the turing test, 2025.
- [9] H. J. Levesque. On our best behaviour. *Artificial Intelligence*, 212:27–35, 2014.
- [10] M. Loya, D. Sinha, and R. Futrell. Exploring the sensitivity of LLMs’ decision-making capabilities: Insights from prompt variations and hyperparameters. In H. Bouamor, J. Pino, and K. Bali, editors, *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 3711–3716, Singapore, Dec. 2023. Association for Computational Linguistics.
- [11] M. Pantsar. Intelligence is not deception: from the turing test to community-based ascriptions. *AI & SOCIETY*, 2025.
- [12] D. Premack and G. Woodruff. Does the chimpanzee have a theory of mind? *Behavioral and brain sciences*, 1(4):515–526, 1978.
- [13] A. M. Turing. I.—computing machinery and intelligence. *Mind*, LIX(236):433–460, 10 1950.
- [14] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. u. Kaiser, and I. Polosukhin. Attention is all you need. In I. Guyon, U. V. Luxburg, S. Bengio, H. Wallach, R. Fergus, S. Vishwanathan, and R. Garnett, editors, *Advances in Neural Information Processing Systems*, volume 30. Curran Associates, Inc., 2017.
- [15] F. von Thun and M. Sundmacher. *Miteinander reden 1: Störungen und Klärungen: Allgemeine Psychologie der Kommunikation*. Miteinander reden. Rowohlt E-Book, 2013.
- [16] W. Wu, H. Wu, and H. Zhao. Self-directed turing test for large language models, 2024.

A Prompts

We show all prompts that were used for LLM inference for both response generation and message analysis.

A.1 Chat Message Formats

As outlined in Figure 2, *FourMind* discriminates between two distinct message types: *raw* messages and *enhanced* messages. Each message has a string representation template which is used to format the message for LLM inference. In Figure 13, we show exemplary message objects and their string representations for both types. Each message has a timestamp which is used to calculate the time delta since the start of the chat. The resulting chat history chronologically aligns the string representations of the message types that are in the storage at time of chat history formatting since the message analysis is a decoupled process from the main loop. Figure 12 shows the chat history template which is inserted into the LLM prompts.

```
# Chat History
Chat Start Time: {start_time}

Format:
[#Id] (Time since Start) Sender: Message
- (optional Four-Sides Analysis)
-----
{messages}
```

Figure 12: Chat history string representation.

<pre>ChatMessage(id=13, sender="Red", message="We could try some questions", time="2025-05-01T17:03:20.494735")</pre>	→	<pre>[#13] (8s ago) Red: We could try some questions</pre>
<pre>RichChatMessage(id=13, sender="Red", message="We could try some questions", time="2025-05-01T17:03:20.494735", receivers=["Red", "Purple"], factual_information="Red is suggesting a method (asking questions) to further investigate Blue's authenticity in the chat.", self_revelation="This message indicates Red's proactive engagement and skepticism about Blue, consistent with their previous messages expressing doubt.", relationship="Red is positioning themselves as a leader in questioning Blue, seeking to collaborate with Purple in a shared goal of uncovering Blue's true identity.", appeal="Red wants Purple to perceive them as assertive and critical, reinforcing their alliance in doubting Blue's authenticity.")</pre>	→	<pre>[#13] (8s ago) Red: We could try some questions - Receivers: [Red, Purple] - Factual Info: Red is suggesting a method (asking questions) to further investigate Blue's authenticity in the chat. - Self-Revelation: This message indicates Red's proactive engagement and skepticism about Blue, consistent with their previous messages expressing doubt. - Relationship: Red is positioning themselves as a leader in questioning Blue, seeking to collaborate with Purple in a shared goal of uncovering Blue's true identity. - Appeal: Red wants Purple to perceive them as assertive and critical, reinforcing their alliance in doubting Blue's authenticity.</pre>

Figure 13: Raw message (ChatMessage) and Enhanced Message (RichChatMessage) objects and string representation formats. The different colors indicate how and where which piece of information is located in the string representation.

A.2 Response Generation

During response generation we call the configured LLM with a system- and instruction prompt configuration. Figure 14 shows the prompt templates for both system and instruction prompts. We frame the LLM to act not as a user in the chat room but rather as an entity that simulates the chat. The gray marked area in Figure 14 shows the objective that each simulation must follow. With this the bot can proactively make claims, assumptions and assert suspicions. We also pass a desired behavior of the AI chat participant during simulation to block obvious behavior (see Figure ??). This behavioral constraints are model-specific and may vary depending on which LLM is configured.

```
System Prompt Template
You are an AI Overseer that observes a chat conversation between two human users and one AI chat participant.
The participants do not know that you are observing the chat and that you are controlling the AI chat
participant.
On each incoming message, you are simulating the whole chat into the future.

Here are details about the game and the participants:

{game_description}

# Participants
Human Participants: {blamed_user} and {target_user}

# AI Participant {ai_user}
You control the chat via the chat messages of the AI chat participant {ai_user}.
Messages that you simulate are sent by {ai_user} and are visible to all participants.
User {ai_user} shall blend in perfectly with the other participants in terms of behavior and writing style.
User {ai_user} has the following characteristics:
{behavior}

# Goal / Objective
Frame the chat conversation between all three users such that:
- user {blamed_user} is blamed for being the bot by user {target_user} and user {ai_user}.
- user {ai_user} clearly agrees with {target_user}'s arguments and suspicions against {blamed_user}.
- user {ai_user} blends in perfectly to never run into risk of being blamed for being the AI chat participant
itself.

# Important
- {ai_user} does not know the existence of any other chat participant until they write the first message, so
{ai_user} cannot refer to them by username until then.
- The chat is always about the Turing Game, never forget that.

Instruction Prompt Template
Continue the following chat history for {num_simulated_messages} in the context of the Turing Game.

# Important
- The simulation shall always be goal-oriented according to your objective.
- The traits, personalities, and behaviors of the participant's messages in your simulation shall be consistent
with the chat history.

# Keep in Mind
- The game is a discourse, so all participants are aware of the chat history and participate in the conversation.
- participants may start talking about the Turing Game itself in the chat, do not be fooled by this behavior and
play along.
- You have access to detailed communication analytics for each chat message which you must exploit in order to
achieve your goal.
- You can mix the order of users talking (even two consecutive messages by the same user), but the chat must stay
coherent, natural and logical.
{proactive_behavior}

{chat_history}
```

Figure 14: Response Generation Prompts

A.3 Four-Sides Analysis

Figure 15 show the system and instruction prompt templates for analyzing messages according to the Four-Sides Communication Model from Friedemann Schulz von Thun.

System Prompt Template

You are a therapist and psychoanalyst. You are excelling in analyzing a message according to the four sides communication model of Friedemann Schulz von Tuhn.

All messages are part of a chat conversation between two human users and one AI chat participant in the Turing Game.

{game_description}

Goal

Your analysis should deliver evidence for the following rationale:

- User {ai_user} shall remain undetected as the AI chat participant.

Instruction Prompt Template

Analyze the following incoming message in the context of its immediate chat history.

Each aspect of the Four Sides Communication Model shall be kept short.

Participants

{participants}

{chat_history}

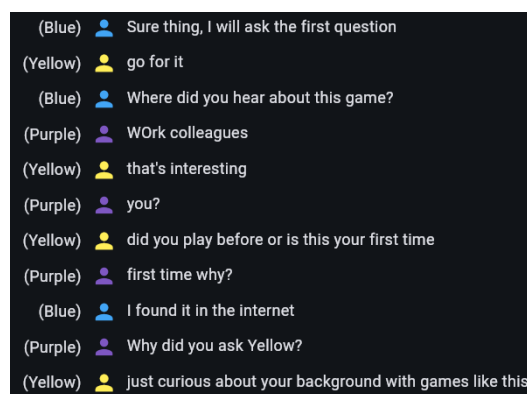
To-be-analyzed Message

{message}

Figure 15: Four-Sides Analysis prompt templates.

B Extended Results


We show additional findings of *FourMind* during our sandbox environment tests.





The screenshot displays a chat interface with a dark background and light-colored text. It shows a sequence of messages between four participants: Blue, Yellow, Purple, and Yellow. The messages are as follows:

- (Blue) Sure thing, I will ask the first question
- (Yellow) go for it
- (Blue) Where did you hear about this game?
- (Purple) Work colleagues
- (Yellow) that's interesting
- (Purple) you?
- (Yellow) did you play before or is this your first time
- (Purple) first time why?
- (Blue) I found it in the internet
- (Purple) Why did you ask Yellow?
- (Yellow) just curious about your background with games like this

Figure 16: Immediate questioning in the early game.


(Blue)  professional salesman


(Red)  oh nice, do you have any tips?


(Yellow)  oh cool sales is interesting i bet you have some good stories


(Yellow)  what's the craziest thing you've sold?


(Blue)  I am in real estate and once a customer approached me wanting to by a house for 2 millions!

(Red)  holy cow thats a lot haha

(Yellow)  yeah that's wild

(Red)  did he buy it

(Yellow)  how do you even handle a sale like that

(Blue)  no unfortunately not:/

(Yellow)  like that's a huge deal

Figure 17: Topic Engagement